SAC network: CNN for the time series, self-attention across EV slots, N + 1 power outputs

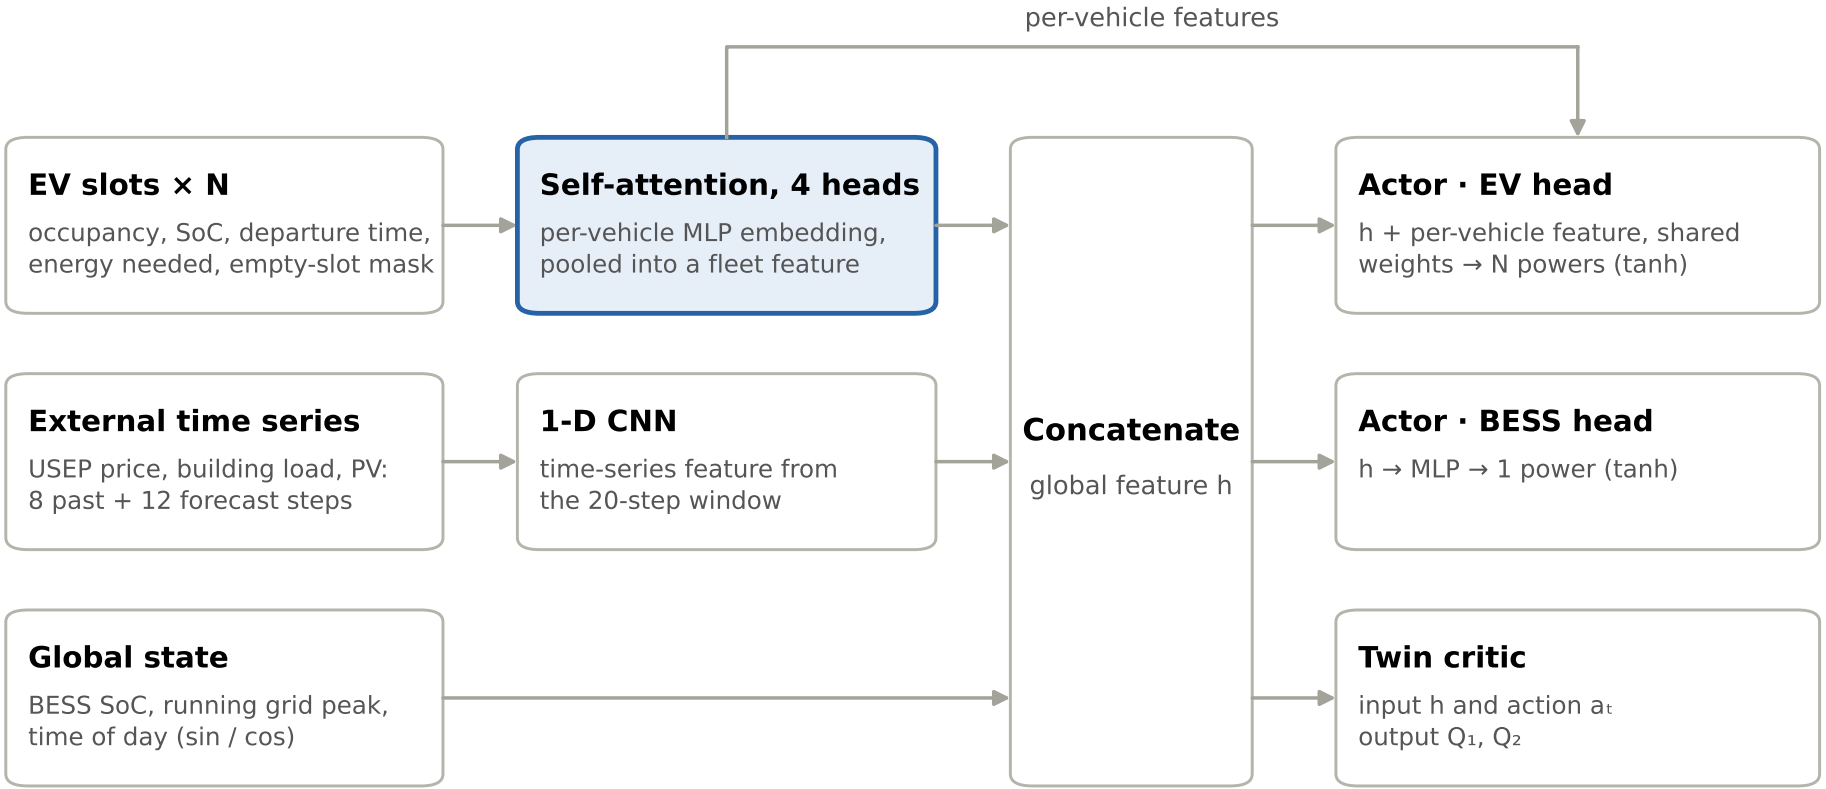

Training: the encoders are updated only by the critic loss; the actor's gradient is detached at the encoder outputs.

Deployment: each action is mapped to its device's feasible power range (0 → 0 kW when idling is feasible) and projected by the environment.